

King County Sales Assessment

```
In [1]: # Import relevant modules
import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
import seaborn as sns
import statsmodels.api as sm
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.datasets import make_regression
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_validate, ShuffleSplit
import sklearn.metrics as metrics
from random import gauss
from mpl_toolkits.mplot3d import Axes3D
from scipy import stats as stats


%matplotlib inline
```

```
In [2]: # Import data
df = pd.read_csv('data/kc_house_data.csv')
df.head()
```

Out[2]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	NaN
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	NO
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	NO
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	NO
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	NO

5 rows × 21 columns



In [3]: `df.describe()`

Out[3]:


	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	
count	2.159700e+04	2.159700e+04	21597.000000	21597.000000	21597.000000	2.159700e+04	21597
mean	4.580474e+09	5.402966e+05	3.373200	2.115826	2080.321850	1.509941e+04	1
std	2.876736e+09	3.673681e+05	0.926299	0.768984	918.106125	4.141264e+04	0
min	1.000102e+06	7.800000e+04	1.000000	0.500000	370.000000	5.200000e+02	1
25%	2.123049e+09	3.220000e+05	3.000000	1.750000	1430.000000	5.040000e+03	1
50%	3.904930e+09	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03	1
75%	7.308900e+09	6.450000e+05	4.000000	2.500000	2550.000000	1.068500e+04	2
max	9.900000e+09	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06	3

In [4]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     21597 non-null  int64
1   date                   21597 non-null  object
2   price                  21597 non-null  float64
3   bedrooms               21597 non-null  int64
4   bathrooms              21597 non-null  float64
5   sqft_living            21597 non-null  int64
6   sqft_lot               21597 non-null  int64
7   floors                 21597 non-null  float64
8   waterfront             19221 non-null  object
9   view                   21534 non-null  object
10  condition              21597 non-null  object
11  grade                  21597 non-null  object
12  sqft_above             21597 non-null  int64
13  sqft_basement          21597 non-null  object
14  yr_built               21597 non-null  int64
15  yr_renovated           17755 non-null  float64
16  zipcode                 21597 non-null  int64
17  lat                    21597 non-null  float64
18  long                   21597 non-null  float64
19  sqft_living15          21597 non-null  int64
20  sqft_lot15             21597 non-null  int64
dtypes: float64(6), int64(9), object(6)
memory usage: 3.5+ MB
```

Baseline Models

```
In [5]: # Lets throw up a heat map to see our simple correlation matrix
plt.figure(figsize=(14, 6))
sns.heatmap(df.corr(), annot=True)
plt.title('Correlation Matrix for Features')
plt.show()
```



```
In [6]: # putting price in the logarithmic scale raises r squared coefficients later on,
df['l_price'] = np.log(df['price'])
```

```
In [7]: # log scaling the living space square footage does not change the r squared much,
df['sqft_living_trans'] = np.log(df['sqft_living'])
```

```
In [8]: # Grade is ordinal, so here we convert the column values to a 1-10 scale, rather
df['grade_final'] = df.grade.map(lambda x: int(x[0]))
# we dont use Grade in the final model, but I feel it is pertinent to show an ex
```

```
In [9]: # Drop NA values to Least common denominator to allow for most comparisons across
df_clean = df.dropna(axis = 0, how = 'any')
```

In [10]: `df_clean.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 15762 entries, 1 to 21596
Data columns (total 24 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    15762 non-null  int64
1   date                  15762 non-null  object
2   price                 15762 non-null  float64
3   bedrooms              15762 non-null  int64
4   bathrooms             15762 non-null  float64
5   sqft_living           15762 non-null  int64
6   sqft_lot              15762 non-null  int64
7   floors                15762 non-null  float64
8   waterfront            15762 non-null  object
9   view                  15762 non-null  object
10  condition              15762 non-null  object
11  grade                  15762 non-null  object
12  sqft_above             15762 non-null  int64
13  sqft_basement          15762 non-null  object
14  yr_built                15762 non-null  int64
15  yr_renovated           15762 non-null  float64
16  zipcode                15762 non-null  int64
17  lat                    15762 non-null  float64
18  long                   15762 non-null  float64
19  sqft_living15          15762 non-null  int64
20  sqft_lot15             15762 non-null  int64
21  l_price                15762 non-null  float64
22  sqft_living_trans      15762 non-null  float64
23  grade_final            15762 non-null  int64
dtypes: float64(8), int64(10), object(6)
memory usage: 3.0+ MB
```

In [11]: *# Create dummies for categorical variables*

```
df_dum = pd.get_dummies(df_clean, columns = ['waterfront', 'zipcode'], drop_first
```

In [12]: df_dum.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 15762 entries, 1 to 21596
Data columns (total 92 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    15762 non-null  int64
1   date                                 15762 non-null  object
2   price                                15762 non-null  float64
3   bedrooms                             15762 non-null  int64
4   bathrooms                             15762 non-null  float64
5   sqft_living                           15762 non-null  int64
6   sqft_lot                              15762 non-null  int64
7   floors                                15762 non-null  float64
8   view                                  15762 non-null  object
9   condition                             15762 non-null  object
10  grade                                 15762 non-null  object
11  sqft_above                            15762 non-null  int64
12  sqft_basement                         15762 non-null  object
13  yr_built                              15762 non-null  int64
14  yr_renovated                          15762 non-null  float64
15  lat                                    15762 non-null  float64
16  long                                   15762 non-null  float64
17  sqft_living15                         15762 non-null  int64
18  sqft_lot15                            15762 non-null  int64
19  l_price                               15762 non-null  float64
20  sqft_living_trans                     15762 non-null  float64
21  grade_final                           15762 non-null  int64
22  waterfront_YES                        15762 non-null  uint8
23  zipcode_98002                         15762 non-null  uint8
24  zipcode_98003                         15762 non-null  uint8
25  zipcode_98004                         15762 non-null  uint8
26  zipcode_98005                         15762 non-null  uint8
27  zipcode_98006                         15762 non-null  uint8
28  zipcode_98007                         15762 non-null  uint8
29  zipcode_98008                         15762 non-null  uint8
30  zipcode_98010                         15762 non-null  uint8
31  zipcode_98011                         15762 non-null  uint8
32  zipcode_98014                         15762 non-null  uint8
33  zipcode_98019                         15762 non-null  uint8
34  zipcode_98022                         15762 non-null  uint8
35  zipcode_98023                         15762 non-null  uint8
36  zipcode_98024                         15762 non-null  uint8
37  zipcode_98027                         15762 non-null  uint8
38  zipcode_98028                         15762 non-null  uint8
39  zipcode_98029                         15762 non-null  uint8
40  zipcode_98030                         15762 non-null  uint8
41  zipcode_98031                         15762 non-null  uint8
42  zipcode_98032                         15762 non-null  uint8
43  zipcode_98033                         15762 non-null  uint8
44  zipcode_98034                         15762 non-null  uint8
45  zipcode_98038                         15762 non-null  uint8
46  zipcode_98039                         15762 non-null  uint8
47  zipcode_98040                         15762 non-null  uint8
48  zipcode_98042                         15762 non-null  uint8
49  zipcode_98045                         15762 non-null  uint8
```

```
50 zipcode_98052      15762 non-null uint8
51 zipcode_98053      15762 non-null uint8
52 zipcode_98055      15762 non-null uint8
53 zipcode_98056      15762 non-null uint8
54 zipcode_98058      15762 non-null uint8
55 zipcode_98059      15762 non-null uint8
56 zipcode_98065      15762 non-null uint8
57 zipcode_98070      15762 non-null uint8
58 zipcode_98072      15762 non-null uint8
59 zipcode_98074      15762 non-null uint8
60 zipcode_98075      15762 non-null uint8
61 zipcode_98077      15762 non-null uint8
62 zipcode_98092      15762 non-null uint8
63 zipcode_98102      15762 non-null uint8
64 zipcode_98103      15762 non-null uint8
65 zipcode_98105      15762 non-null uint8
66 zipcode_98106      15762 non-null uint8
67 zipcode_98107      15762 non-null uint8
68 zipcode_98108      15762 non-null uint8
69 zipcode_98109      15762 non-null uint8
70 zipcode_98112      15762 non-null uint8
71 zipcode_98115      15762 non-null uint8
72 zipcode_98116      15762 non-null uint8
73 zipcode_98117      15762 non-null uint8
74 zipcode_98118      15762 non-null uint8
75 zipcode_98119      15762 non-null uint8
76 zipcode_98122      15762 non-null uint8
77 zipcode_98125      15762 non-null uint8
78 zipcode_98126      15762 non-null uint8
79 zipcode_98133      15762 non-null uint8
80 zipcode_98136      15762 non-null uint8
81 zipcode_98144      15762 non-null uint8
82 zipcode_98146      15762 non-null uint8
83 zipcode_98148      15762 non-null uint8
84 zipcode_98155      15762 non-null uint8
85 zipcode_98166      15762 non-null uint8
86 zipcode_98168      15762 non-null uint8
87 zipcode_98177      15762 non-null uint8
88 zipcode_98178      15762 non-null uint8
89 zipcode_98188      15762 non-null uint8
90 zipcode_98198      15762 non-null uint8
91 zipcode_98199      15762 non-null uint8
dtypes: float64(8), int64(9), object(5), uint8(70)
memory usage: 3.8+ MB
```

```
In [13]: df_dum.columns
```

```
Out[13]: Index(['id', 'date', 'price', 'bedrooms', 'bathrooms', 'sqft_living',  
               'sqft_lot', 'floors', 'view', 'condition', 'grade', 'sqft_above',  
               'sqft_basement', 'yr_built', 'yr_renovated', 'lat', 'long',  
               'sqft_living15', 'sqft_lot15', 'l_price', 'sqft_living_trans',  
               'grade_final', 'waterfront_YES', 'zipcode_98002', 'zipcode_98003',  
               'zipcode_98004', 'zipcode_98005', 'zipcode_98006', 'zipcode_98007',  
               'zipcode_98008', 'zipcode_98010', 'zipcode_98011', 'zipcode_98014',  
               'zipcode_98019', 'zipcode_98022', 'zipcode_98023', 'zipcode_98024',  
               'zipcode_98027', 'zipcode_98028', 'zipcode_98029', 'zipcode_98030',  
               'zipcode_98031', 'zipcode_98032', 'zipcode_98033', 'zipcode_98034',  
               'zipcode_98038', 'zipcode_98039', 'zipcode_98040', 'zipcode_98042',  
               'zipcode_98045', 'zipcode_98052', 'zipcode_98053', 'zipcode_98055',  
               'zipcode_98056', 'zipcode_98058', 'zipcode_98059', 'zipcode_98065',  
               'zipcode_98070', 'zipcode_98072', 'zipcode_98074', 'zipcode_98075',  
               'zipcode_98077', 'zipcode_98092', 'zipcode_98102', 'zipcode_98103',  
               'zipcode_98105', 'zipcode_98106', 'zipcode_98107', 'zipcode_98108',  
               'zipcode_98109', 'zipcode_98112', 'zipcode_98115', 'zipcode_98116',  
               'zipcode_98117', 'zipcode_98118', 'zipcode_98119', 'zipcode_98122',  
               'zipcode_98125', 'zipcode_98126', 'zipcode_98133', 'zipcode_98136',  
               'zipcode_98144', 'zipcode_98146', 'zipcode_98148', 'zipcode_98155',  
               'zipcode_98166', 'zipcode_98168', 'zipcode_98177', 'zipcode_98178',  
               'zipcode_98188', 'zipcode_98198', 'zipcode_98199'],  
              dtype='object')
```

```
In [14]: # Establish baseline model for strong predictor, ZIP code
X1, y1 = df_dum.drop(['id', 'date', 'price', 'bedrooms', 'bathrooms', 'sqft_living',
                    'sqft_lot', 'floors', 'view', 'condition', 'grade', 'sqft_above',
                    'sqft_basement', 'yr_built', 'yr_renovated', 'lat', 'long',
                    'sqft_living15', 'sqft_lot15', 'l_price', 'grade_final', 'sqft_living_tram',
                    'waterfront_YES'], axis=1), df_dum['l_price']
X1 = sm.add_constant(X1)
baseline = sm.OLS(y1, X1).fit()
baseline.summary()
```

Out[14]: OLS Regression Results

Dep. Variable:	l_price	R-squared:	0.533
Model:	OLS	Adj. R-squared:	0.531
Method:	Least Squares	F-statistic:	259.3
Date:	Fri, 24 Jun 2022	Prob (F-statistic):	0.00
Time:	16:55:52	Log-Likelihood:	-6249.8
No. Observations:	15762	AIC:	1.264e+04
Df Residuals:	15692	BIC:	1.318e+04
Df Model:	69		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	12.4593	0.023	551.867	0.000	12.415	12.504
zipcode_98002	-0.1222	0.037	-3.286	0.001	-0.195	-0.049
zipcode_98003	0.0649	0.033	1.950	0.051	-0.000	0.130
zipcode_98004	1.5554	0.033	47.497	0.000	1.491	1.620
zipcode_98005	1.1061	0.039	28.397	0.000	1.030	1.182
zipcode_98006	1.1063	0.029	37.513	0.000	1.048	1.164
zipcode_98007	0.8330	0.042	19.652	0.000	0.750	0.916
zipcode_98008	0.8220	0.033	24.593	0.000	0.756	0.887
zipcode_98010	0.4028	0.049	8.280	0.000	0.307	0.498
zipcode_98011	0.6168	0.038	16.227	0.000	0.542	0.691
zipcode_98014	0.5051	0.044	11.567	0.000	0.420	0.591
zipcode_98019	0.4711	0.039	12.063	0.000	0.395	0.548
zipcode_98022	0.1400	0.036	3.929	0.000	0.070	0.210
zipcode_98023	0.0598	0.029	2.037	0.042	0.002	0.117
zipcode_98024	0.6744	0.052	12.859	0.000	0.572	0.777
zipcode_98027	0.7855	0.031	25.678	0.000	0.726	0.845
zipcode_98028	0.5432	0.034	16.191	0.000	0.477	0.609
zipcode_98029	0.8183	0.032	25.476	0.000	0.755	0.881
zipcode_98030	0.1299	0.035	3.737	0.000	0.062	0.198

zipcode_98031	0.1292	0.034	3.809	0.000	0.063	0.196
zipcode_98032	-0.0612	0.043	-1.433	0.152	-0.145	0.023
zipcode_98033	1.0219	0.030	33.697	0.000	0.962	1.081
zipcode_98034	0.6184	0.029	21.364	0.000	0.562	0.675
zipcode_98038	0.2934	0.028	10.326	0.000	0.238	0.349
zipcode_98039	2.0415	0.064	31.805	0.000	1.916	2.167
zipcode_98040	1.4342	0.034	42.407	0.000	1.368	1.500
zipcode_98042	0.1538	0.029	5.343	0.000	0.097	0.210
zipcode_98045	0.4572	0.036	12.624	0.000	0.386	0.528
zipcode_98052	0.8831	0.029	30.801	0.000	0.827	0.939
zipcode_98053	0.8830	0.031	28.643	0.000	0.823	0.943
zipcode_98055	0.1197	0.034	3.511	0.000	0.053	0.187
zipcode_98056	0.4033	0.031	13.164	0.000	0.343	0.463
zipcode_98058	0.2606	0.030	8.640	0.000	0.201	0.320
zipcode_98059	0.5620	0.030	18.673	0.000	0.503	0.621
zipcode_98065	0.6681	0.033	20.212	0.000	0.603	0.733
zipcode_98070	0.5498	0.044	12.388	0.000	0.463	0.637
zipcode_98072	0.7376	0.034	21.752	0.000	0.671	0.804
zipcode_98074	0.9405	0.031	30.765	0.000	0.881	1.000
zipcode_98075	1.0708	0.031	34.043	0.000	1.009	1.132
zipcode_98077	0.8938	0.037	24.342	0.000	0.822	0.966
zipcode_98092	0.2122	0.032	6.680	0.000	0.150	0.275
zipcode_98102	1.1390	0.050	22.737	0.000	1.041	1.237
zipcode_98103	0.7499	0.029	26.260	0.000	0.694	0.806
zipcode_98105	1.0995	0.036	30.963	0.000	1.030	1.169
zipcode_98106	0.1653	0.033	5.066	0.000	0.101	0.229
zipcode_98107	0.7734	0.034	23.020	0.000	0.708	0.839
zipcode_98108	0.2875	0.038	7.474	0.000	0.212	0.363
zipcode_98109	1.0809	0.047	22.941	0.000	0.989	1.173
zipcode_98112	1.3419	0.034	39.837	0.000	1.276	1.408
zipcode_98115	0.8248	0.029	28.644	0.000	0.768	0.881
zipcode_98116	0.7987	0.032	24.891	0.000	0.736	0.862
zipcode_98117	0.7549	0.029	26.169	0.000	0.698	0.811
zipcode_98118	0.4030	0.029	13.719	0.000	0.345	0.461
zipcode_98119	1.1250	0.038	29.247	0.000	1.050	1.200
zipcode_98122	0.8088	0.034	24.042	0.000	0.743	0.875
zipcode_98125	0.5227	0.031	16.984	0.000	0.462	0.583

zipcode_98126	0.4448	0.032	13.890	0.000	0.382	0.508
zipcode_98133	0.3769	0.030	12.668	0.000	0.319	0.435
zipcode_98136	0.6987	0.034	20.373	0.000	0.631	0.766
zipcode_98144	0.6993	0.032	22.029	0.000	0.637	0.762
zipcode_98146	0.1927	0.033	5.767	0.000	0.127	0.258
zipcode_98148	0.0519	0.060	0.865	0.387	-0.066	0.170
zipcode_98155	0.3953	0.030	12.990	0.000	0.336	0.455
zipcode_98166	0.4791	0.035	13.759	0.000	0.411	0.547
zipcode_98168	-0.0873	0.034	-2.538	0.011	-0.155	-0.020
zipcode_98177	0.8561	0.035	24.665	0.000	0.788	0.924
zipcode_98178	0.1134	0.035	3.278	0.001	0.046	0.181
zipcode_98188	0.0582	0.043	1.353	0.176	-0.026	0.142
zipcode_98198	0.0662	0.034	1.955	0.051	-0.000	0.133
zipcode_98199	1.0240	0.033	30.904	0.000	0.959	1.089

Omnibus:	1583.149	Durbin-Watson:	1.985
Prob(Omnibus):	0.000	Jarque-Bera (JB):	3910.041
Skew:	0.594	Prob(JB):	0.00
Kurtosis:	5.132	Cond. No.	67.1

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [15]: X2, y2 = df['sqft_living_trans'], df['l_price']
plt.figure(figsize=(9,6))


plt.plot(X2, y2, 'o')

# get m (slope) and b(intercept) of linear regression line
m, b = np.polyfit(X2, y2, 1)

# add linear regression line to scatterplot
plt.plot(X2, m*X2+b, label = 'R2 = 0.455')

plt.title('Living Area Profitability')
plt.xlabel('Living Area (log square feet)')
plt.ylabel('Sale Price (log USD)')
plt.legend()
;
```

Out[15]: ''



Final Model

In order to cut down on unnecessary code as per the guidelines, I omit the model summaries for multiple regressions that included grade and waterfront status in favor of the final model. Adding grade as an ordinal predictor does not increase the accuracy of the model at all, and mildly increases skewness. I do not currently

```
In [16]: # Assign final predictors and dependent variable
X, y = df_dum.drop(['id', 'date', 'price', 'bedrooms', 'bathrooms',
                    'sqft_lot', 'floors', 'view', 'condition', 'grade', 'sqft_above',
                    'sqft_basement', 'yr_built', 'yr_renovated', 'lat', 'long',
                    'sqft_living15', 'sqft_lot15', 'l_price',
                    'sqft_living', 'grade_final'
                    ], axis=1), df_dum['l_price']
```

```
In [17]: X = sm.add_constant(X)
model1 = sm.OLS(y, X).fit()
model1.summary()
```

Out[17]: OLS Regression Results

Dep. Variable:	l_price	R-squared:	0.833
Model:	OLS	Adj. R-squared:	0.832
Method:	Least Squares	F-statistic:	1100.
Date:	Fri, 24 Jun 2022	Prob (F-statistic):	0.00
Time:	16:55:53	Log-Likelihood:	1842.6
No. Observations:	15762	AIC:	-3541.
Df Residuals:	15690	BIC:	-2989.
Df Model:	71		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	7.0687	0.036	193.898	0.000	6.997	7.140

```
In [18]: # Prepare train and test data from X and y variables for cross validation
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_s

linreg = LinearRegression()
linreg.fit(X_train, y_train)

y_hat_train = linreg.predict(X_train)
y_hat_test = linreg.predict(X_test)
```

```
In [19]: # Inspect residuals
train_residuals = y_hat_train - y_train
test_residuals = y_hat_test - y_test
print(train_residuals)
print(test_residuals)
```

18060 0.091455
4178 -0.018608
13219 0.082941
2839 -0.266770
5520 0.256757
...
7129 0.031156
18336 0.246032
7384 -0.242698
1159 -0.580698
9955 0.115760
Name: l_price, Length: 12609, dtype: float64
8446 0.379067
7473 -0.238393
20534 0.142403
5063 0.147336
17989 -0.215431
...
2330 -0.024277
10260 0.106158
16786 0.203038
8159 0.079904
10136 0.255806
Name: l_price, Length: 3153, dtype: float64


```
In [20]: # Determine MSE
mse_train = np.sum((y_train-y_hat_train)**2)/len(y_train)
mse_test = np.sum((y_test-y_hat_test)**2)/len(y_test)
print('Train Mean Squarred Error:', mse_train)
print('Test Mean Squarred Error:', mse_test)
```

Train Mean Squarred Error: 0.04624036654907834
Test Mean Squarred Error: 0.04704959588781619

```
In [21]: # Create a residual plot to explore possible patterns
df_dum["predicted"] = model1.predict(X)
df_dum["residuals"] = model1.resid

plt.figure(figsize=(10, 6))
sns.scatterplot(data=df_dum, x="predicted", y="residuals", alpha = 0.2)
plt.axhline(y=0)
plt.title('Residual Plot')
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
;
```

Out[21]: ''



```
In [22]: splitter = ShuffleSplit(n_splits=3, test_size=0.25, random_state=0)

baseline_scores = cross_validate(estimator=linreg, X=X_train, y=y_train, return_t

print("Train score:      ", baseline_scores["train_score"].mean())
print("Validation score:", baseline_scores["test_score"].mean())
```

Train score: 0.8354967482802792

Validation score: 0.8314829925222368